# Volcano for the Agentic AI Era:

## Unified Scheduling Across Training, Inference, and Agents

**Ricardo Aravena · Rainie Li**

KubeCon + CloudNativeCon North America 2026
Salt Lake City · November 9–12

# Speakers

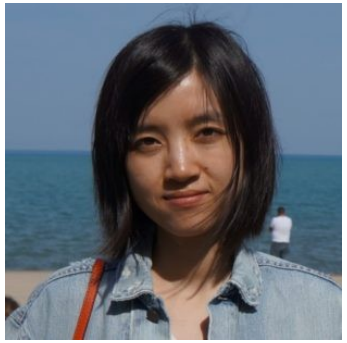


## Ricardo Aravena

AI Infrastructure Lead · Nunchux AI

CNCF TOC Alternate (GB)
Cloud native and AI infrastructure

nunchux.ai · CNCF TOC



## Rainie Li

Pinterest

Data processing infrastructure
Apache YuniKorn PMC member

Apache YuniKorn · rainieli

# 01

# Unified scheduling

Different paths. Shared infrastructure.


**01 Unified scheduling**

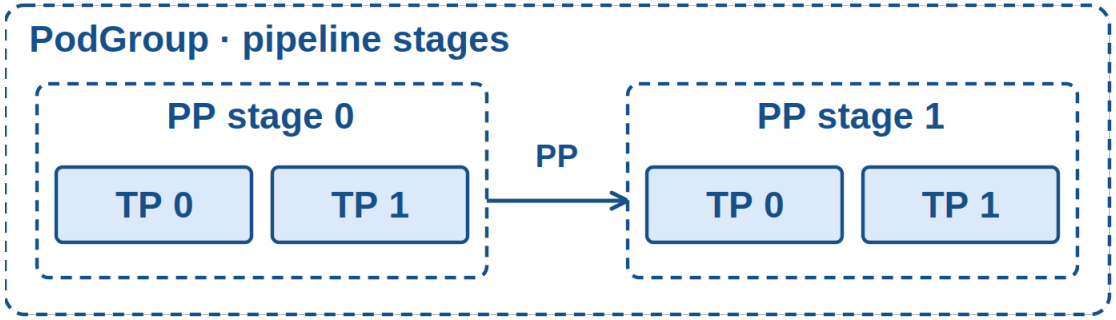
02 Multi-level gang & network topology

03 Ecosystem

04 Benchmark framework & progress

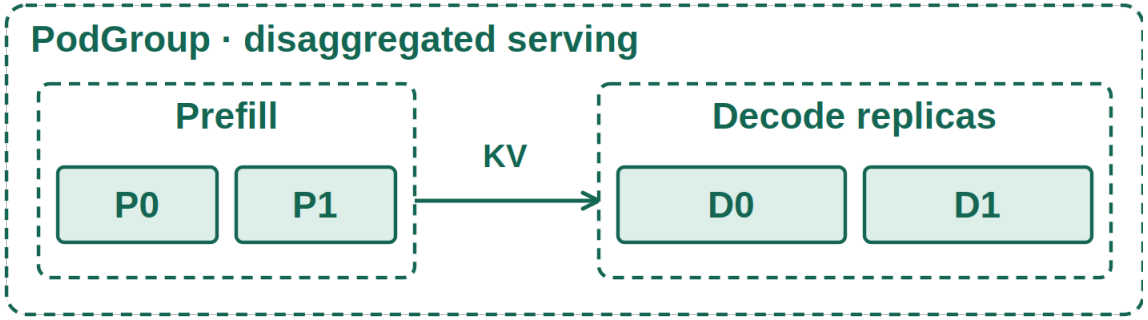
# Different Workloads, Shared Infrastructure

## Training




Gang + locality

## Inference




Gang + locality

## Agents

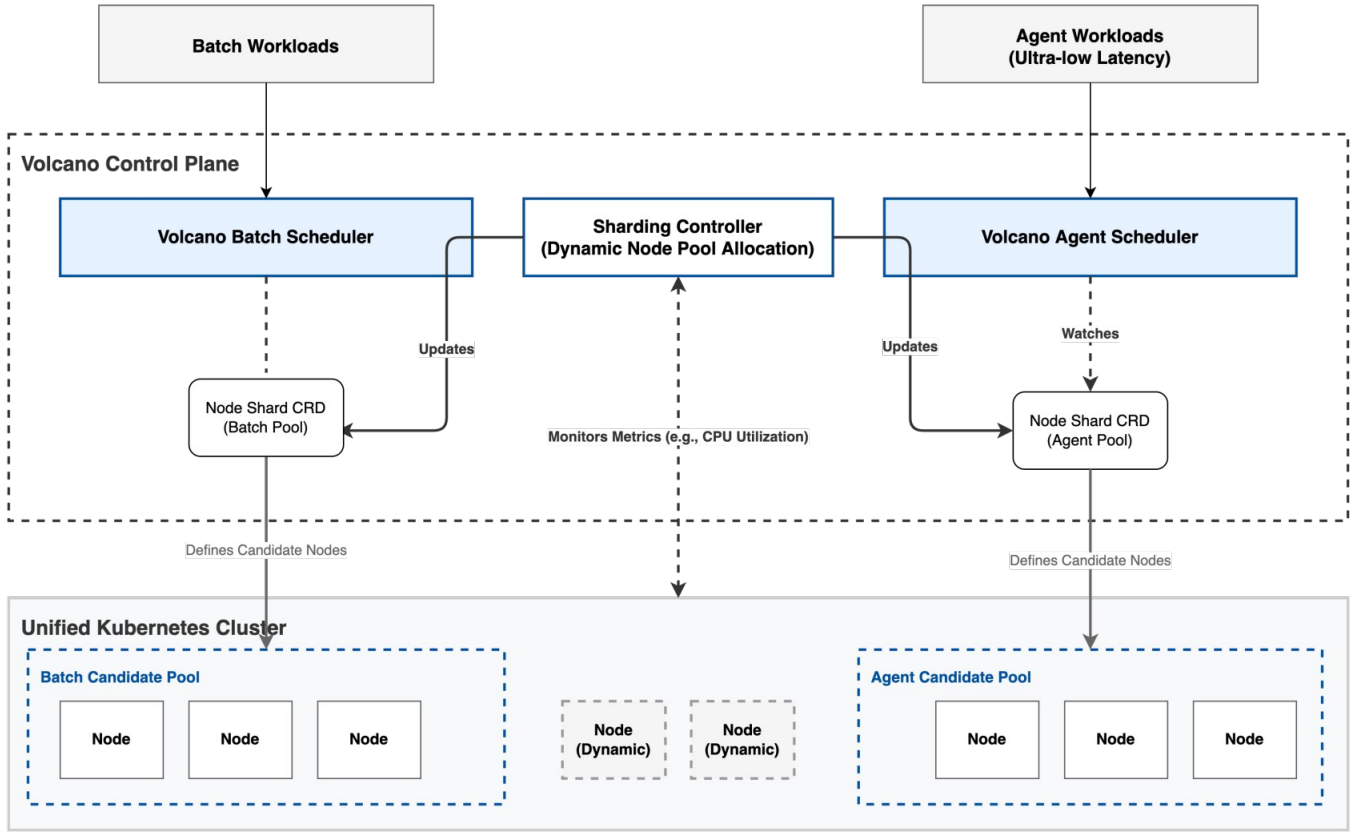
Bursty arrivals



Low startup latency



# A Coordinated Scheduling Framework

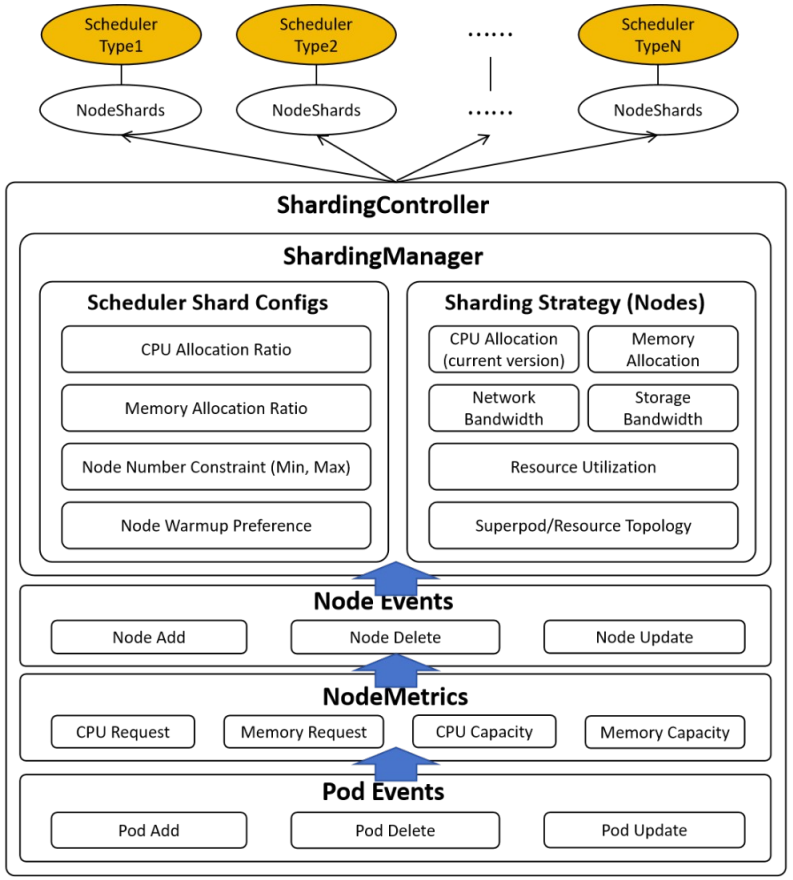



**Batch / fast path**

**Binder checks**

**Shared node pools**

# NodeShard: Desired Nodes and Current Use




**NodeShard · spec / status excerpt**

```
spec:
  nodesDesired:
  - node1
  - node2
status:
  nodesInUse:
  - node0
  - node1
  nodesToRemove:
  - node0
  nodesToAdd:
  - node2
```

# 02

# Multi-level gang & network topology

From available nodes to runnable workloads.

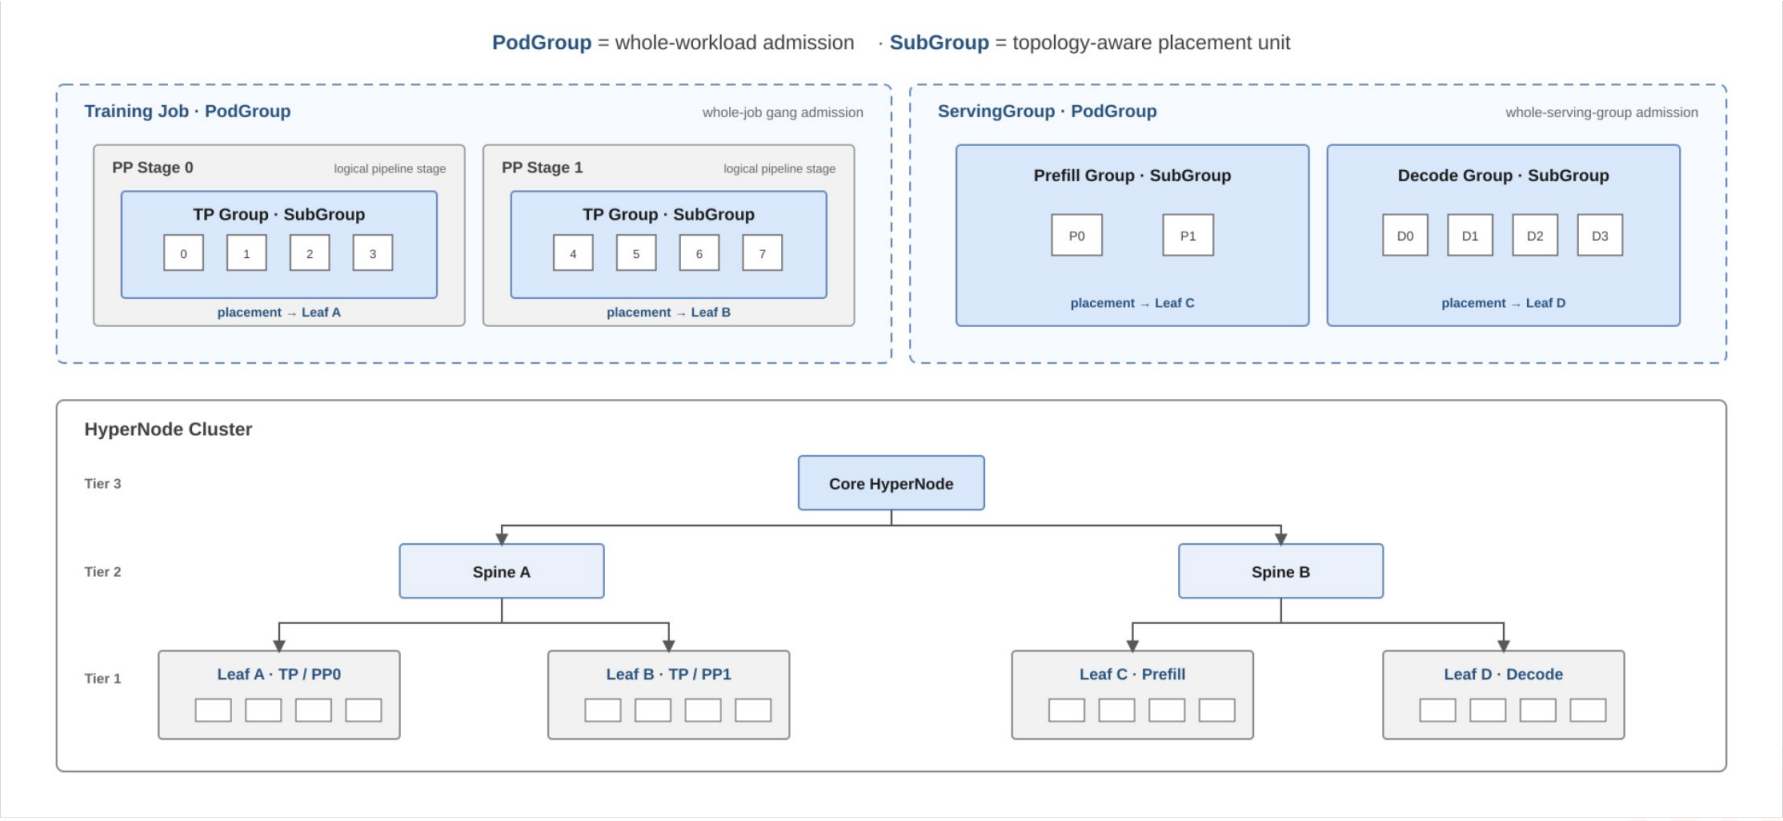

01 Unified scheduling

**02 Multi-level gang & network topology**

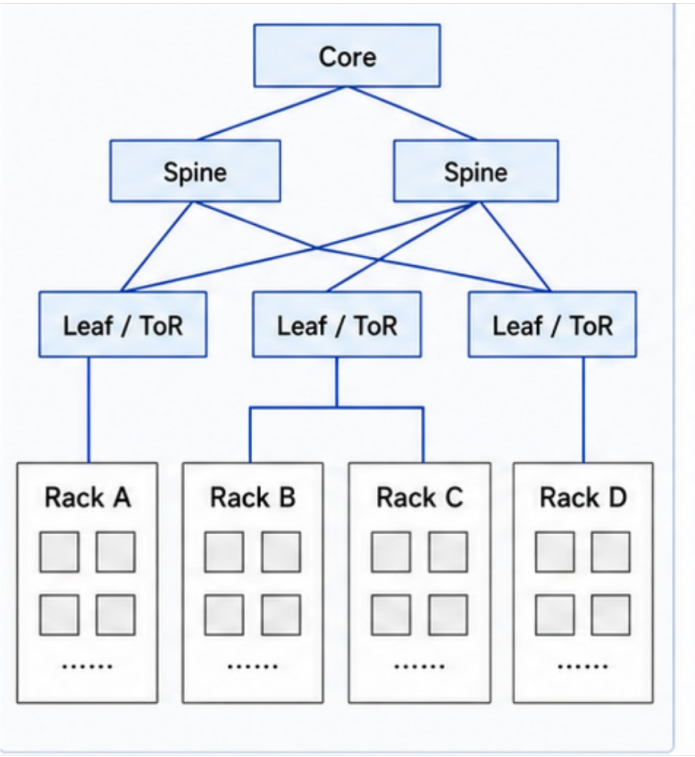
03 Ecosystem

04 Benchmark framework & progress

# PodGroup and SubGroup: Preserve Workload Structure



# HyperNode: From Physical Fabric to Scheduling Domains




**Standalone HyperNode**

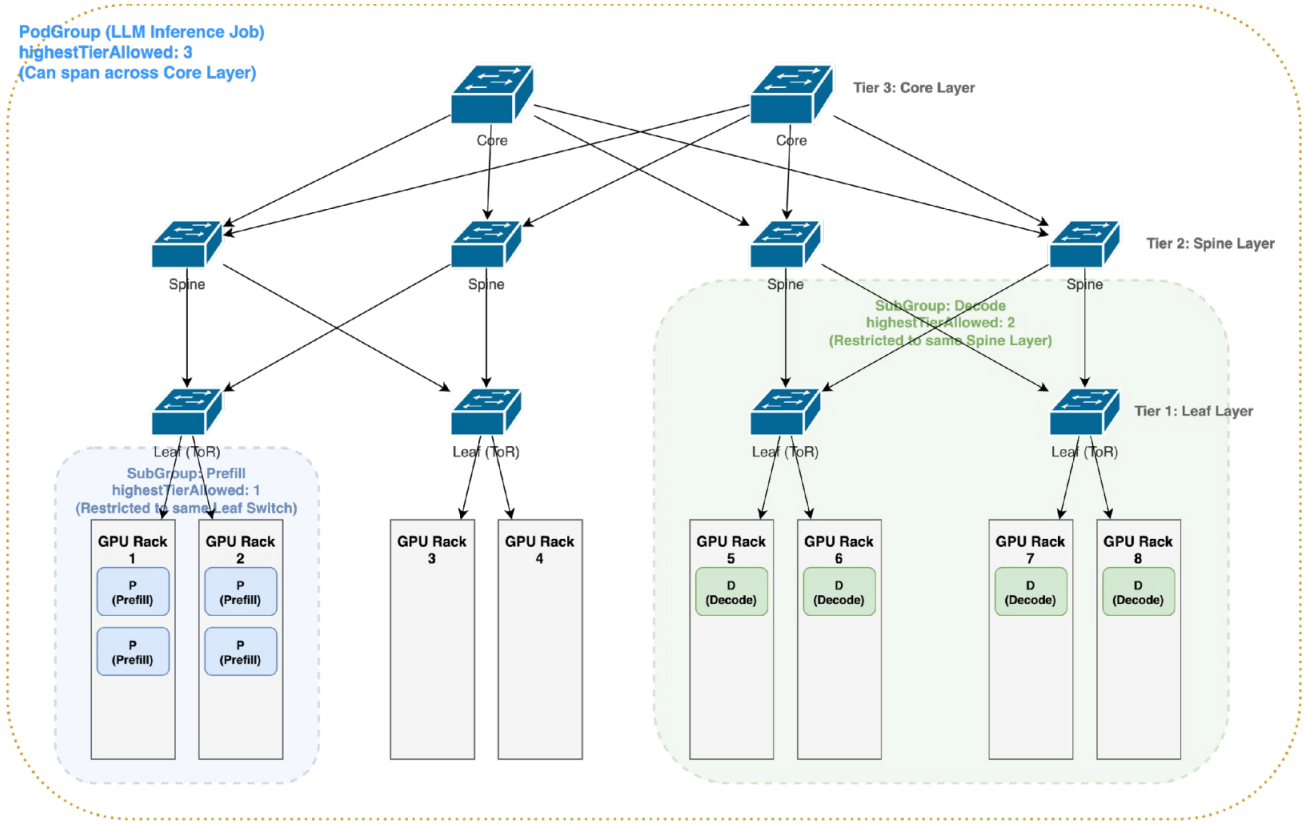
**Parent · HyperNode.spec**

```
tier: 2
tierName: spine
members:
- type: HyperNode
  selector:
    exactMatch:
      name: rack-a
```

**Leaf · HyperNode.spec**

```
tier: 1
tierName: leaf
members:
- type: Node
  selector:
    labelMatch:
      matchLabels:
        rack: rack-a
```

# Place Groups at Different Topology Levels

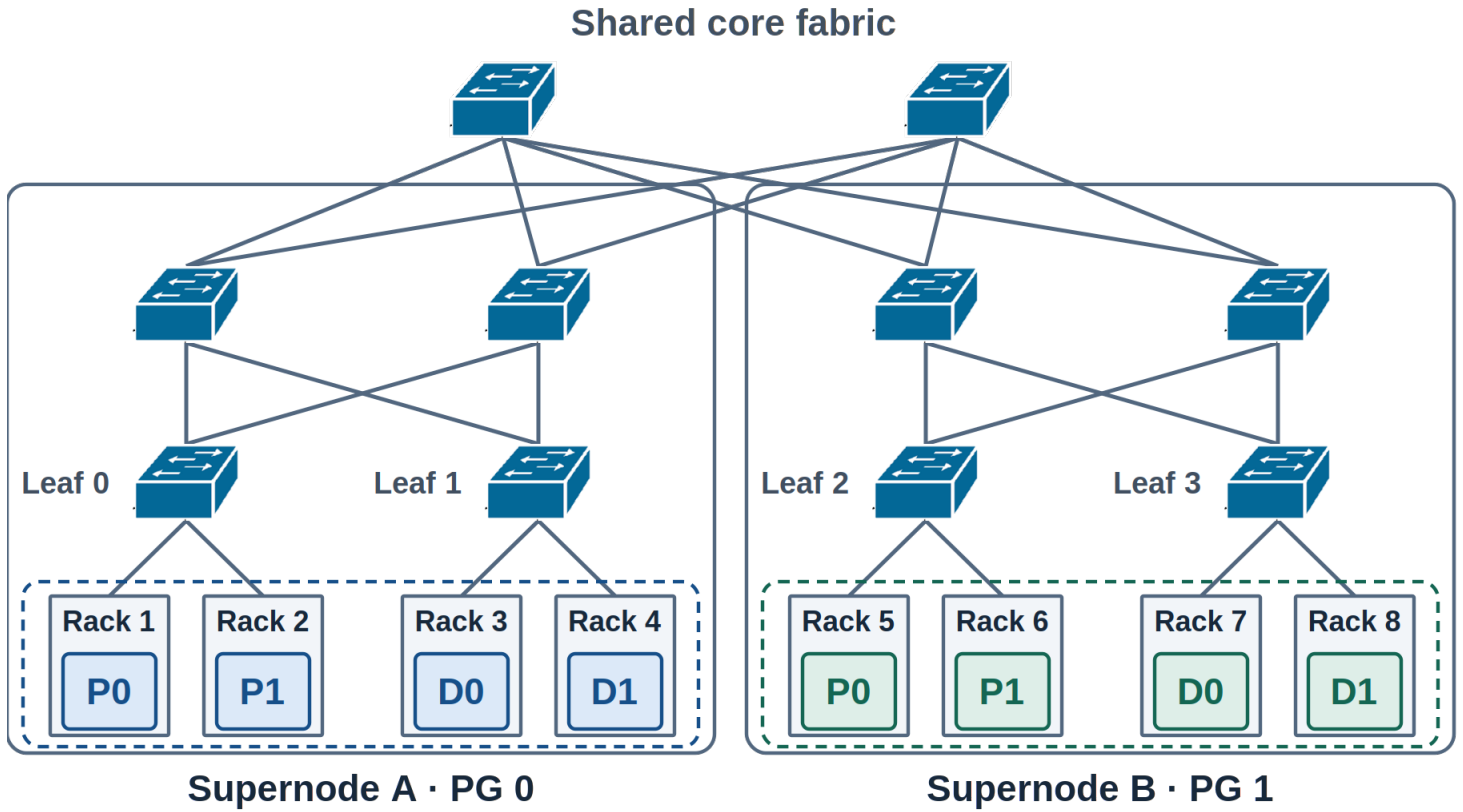



**PodGroup.spec · key fields**

```
minMember: 8
networkTopology:
  highestTierName: core
subGroupPolicy:
- name: prefill
  subGroupSize: 4
  networkTopology:
    highestTierName: leaf
- name: decode
  subGroupSize: 4
  networkTopology:
    highestTierName: spine
```

# Keep Each Instance Together, Place Instances Apart

## Kthena anti-affinity integration: proposed



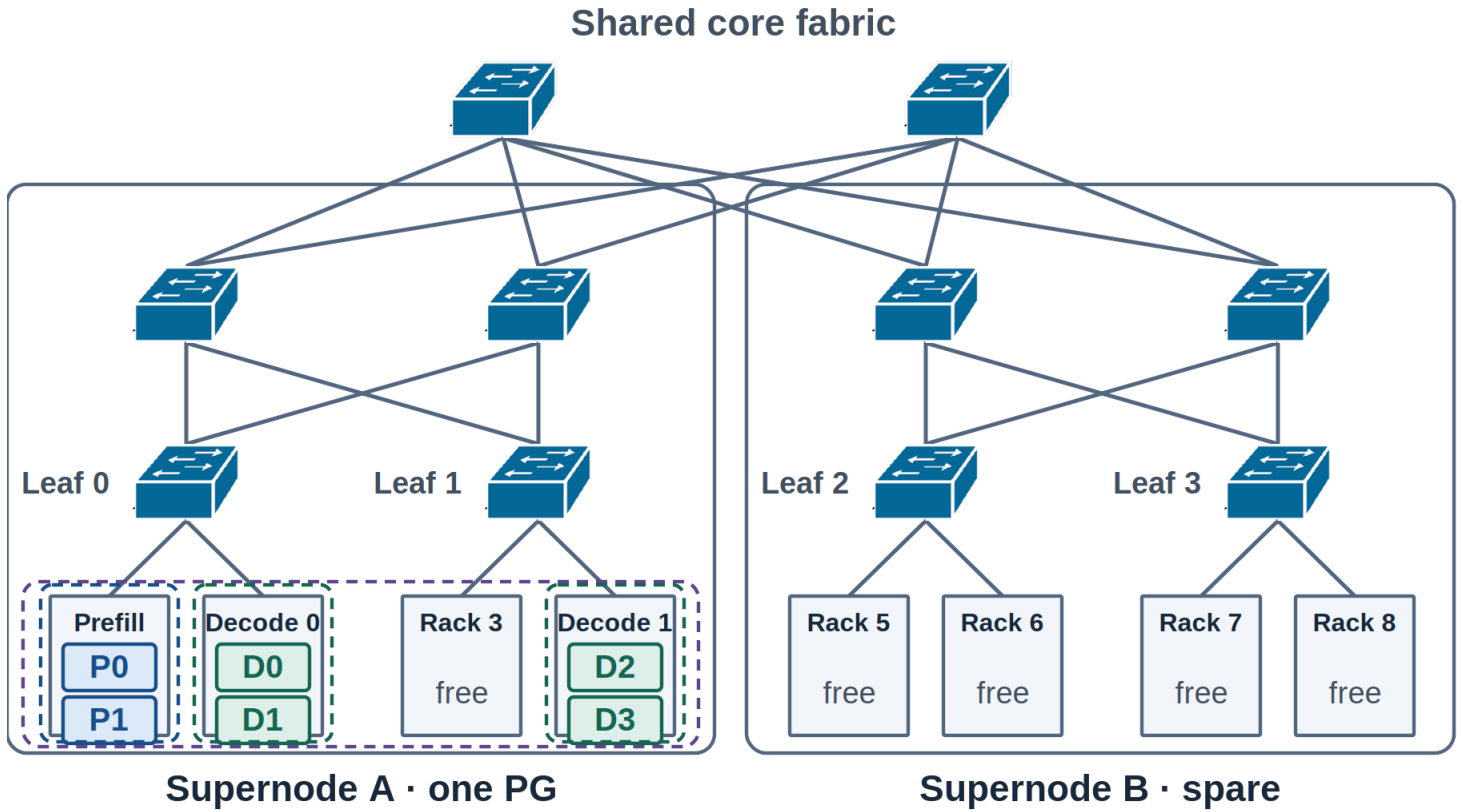

**PodGroup.spec · key fields**

```
networkTopology:
  mode: hard
  highestTierName: supernode
topologyAffinity:
  podGroupAntiAffinity:
    required:
    - podGroupSelector:
        matchLabels:
          app: model-a
      topologyTierName: supernode
```

**PG label: app=model-a**

# Next: Relationships Between SubGroups

## Proposed




### Affinity

Supernode · KV locality

### Anti-affinity

Leaf · replica isolation

# 03

# Ecosystem

Workload controllers carry scheduling intent.

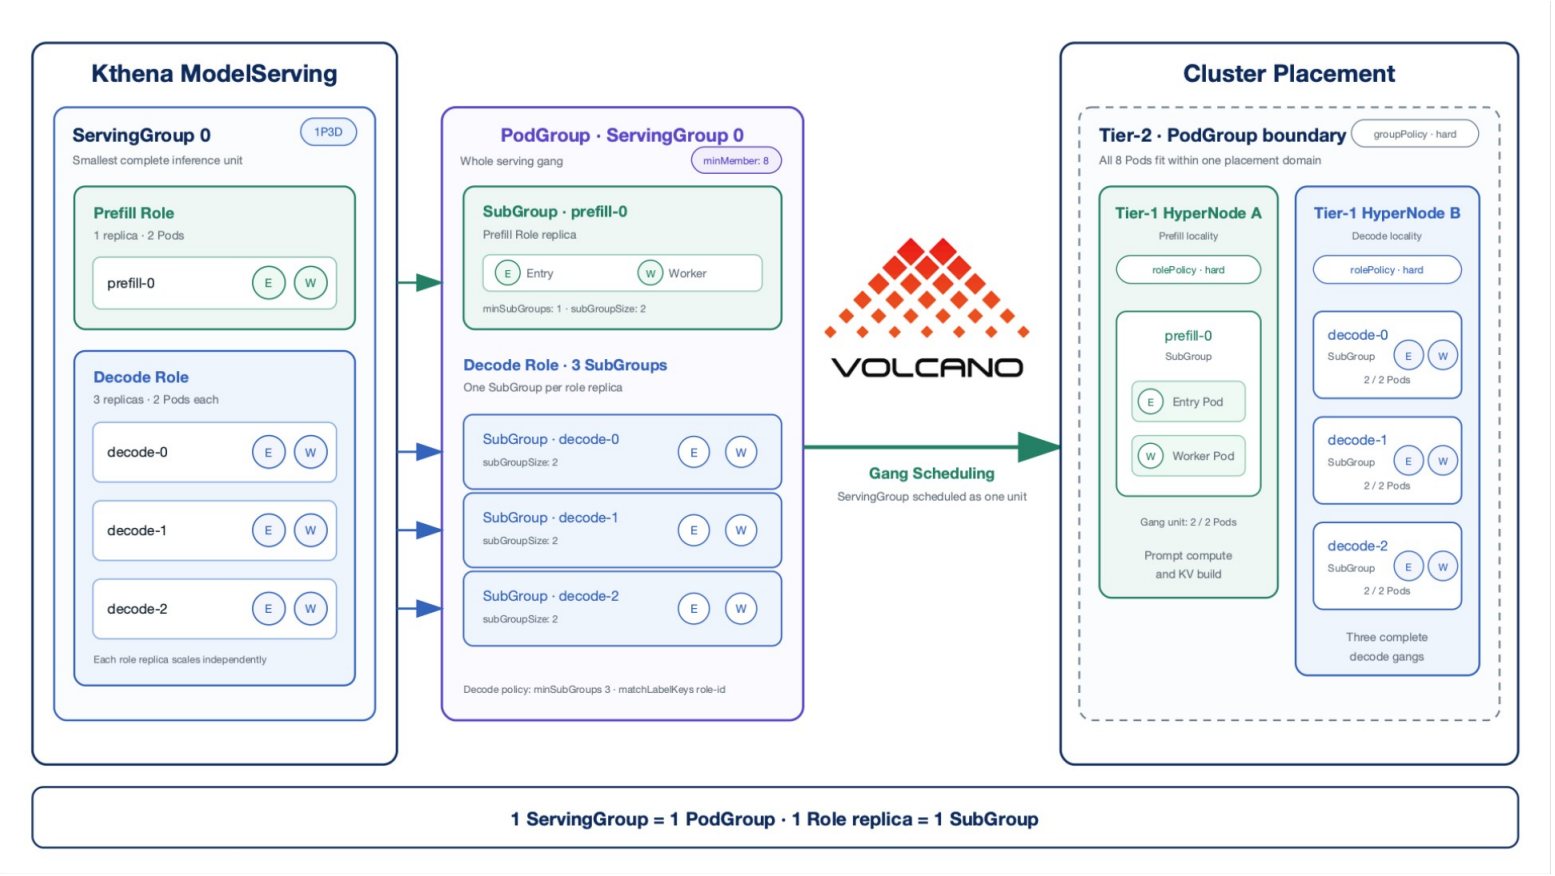

01 Unified scheduling

02 Multi-level gang & network topology

**03 Ecosystem**

04 Benchmark framework & progress

# Kthena: From Serving Roles to Scheduling Groups



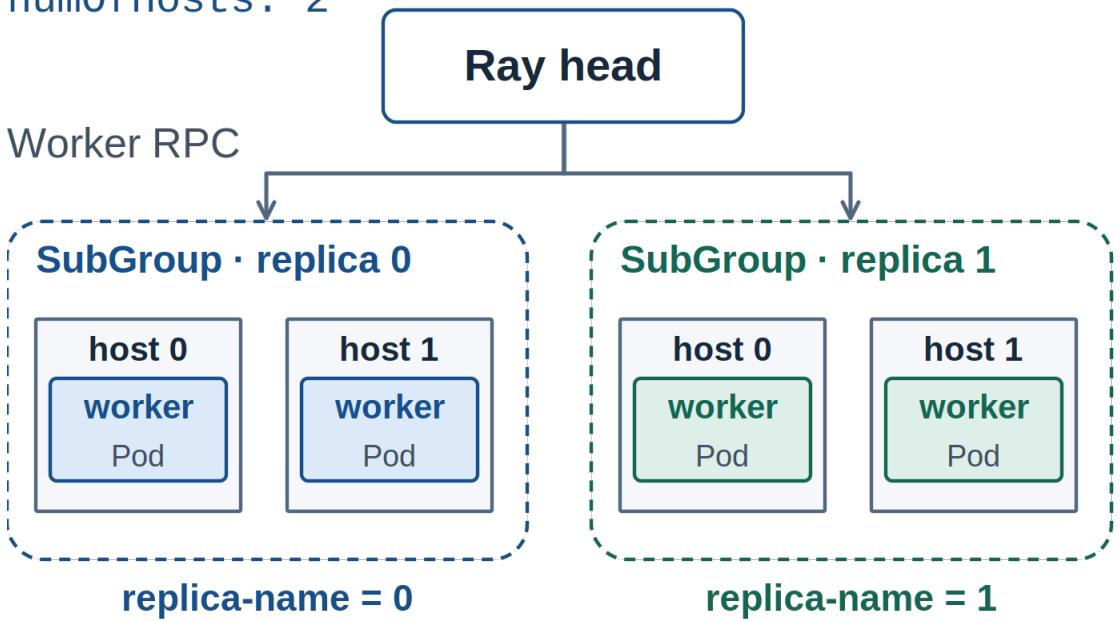

**ModelServing.spec · excerpt**

```
schedulerName: volcano
replicas: 1
template:
  gangPolicy:
    minRoleReplicas:
      prefill: 1
      decode: 3
  roles:
  - name: prefill
    replicas: 1
    workerReplicas: 1
  - name: decode
    replicas: 3
    workerReplicas: 1
```

# KubeRay: A Replica Can Span Multiple Pods

**RayCluster · worker group**

```
minReplicas: 2
numOfHosts: 2
```



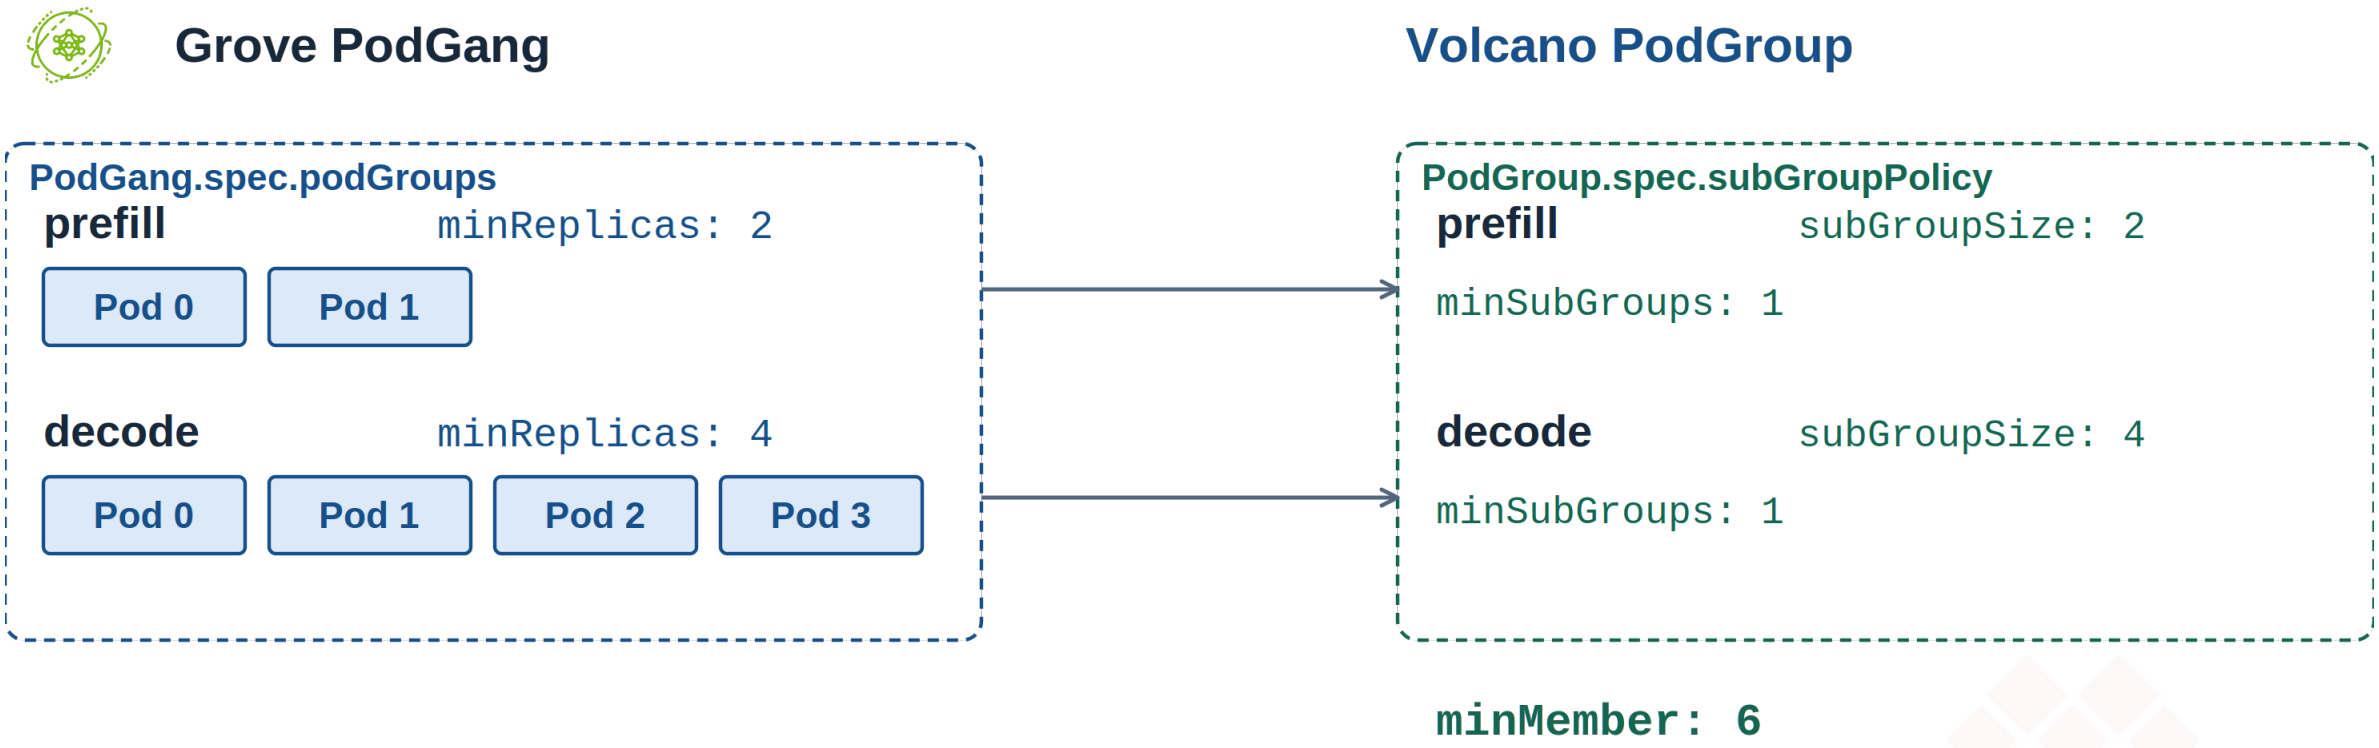

**PodGroup.spec · generated grouping**

```
subGroupPolicy:
- name: worker
  labelSelector:
    matchLabels:
      ray.io/group: worker
  matchLabelKeys:
  - ray.io/worker-group-replica-name
  minSubGroups: 2
  subGroupSize: 2
```

**Gang: merged · Topology: in progress**

# Grove: Carry PodGang Structure into Volcano

## Gang mapping available · topology unsupported



# Next: RBG, Kubeflow, and LWS




**RBG**

In progress

**RoleInstance → SubGroup**

---




**Kubeflow**

Planned

**Training groups + topology**

---




**LWS**

Planned

**Multi-role / multi-group**

# 04

# Benchmark framework & progress

Measure behavior across workload shapes.

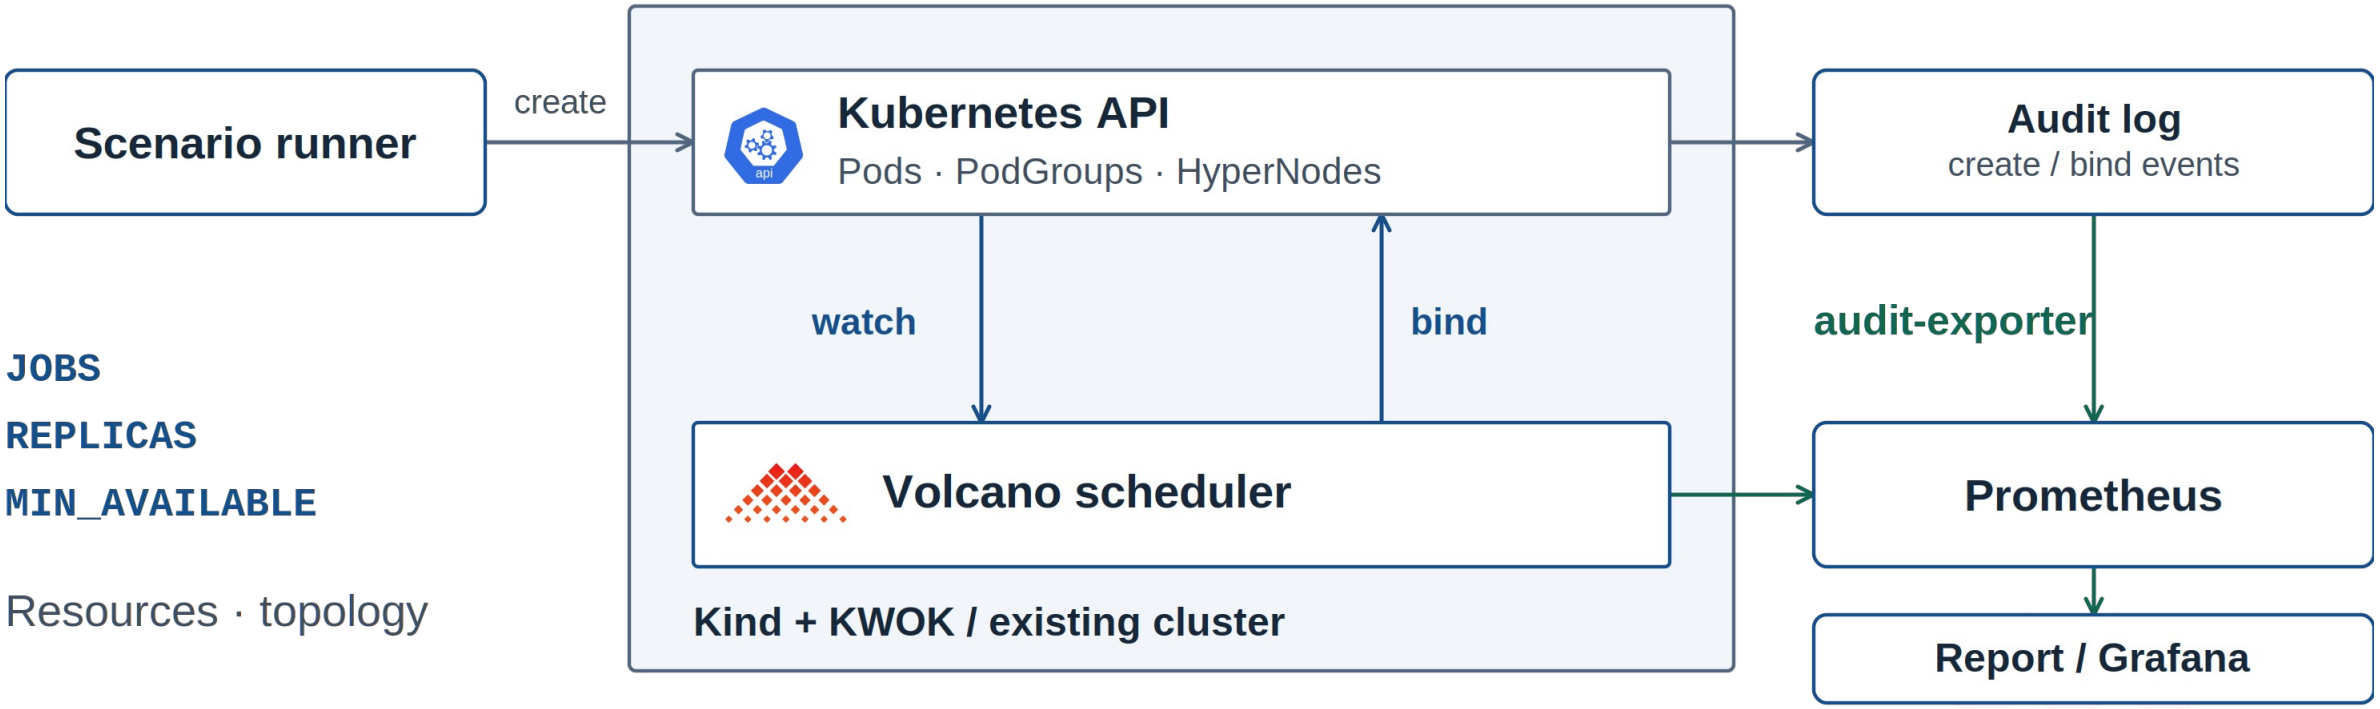

01 Unified scheduling

02 Multi-level gang & network topology

03 Ecosystem

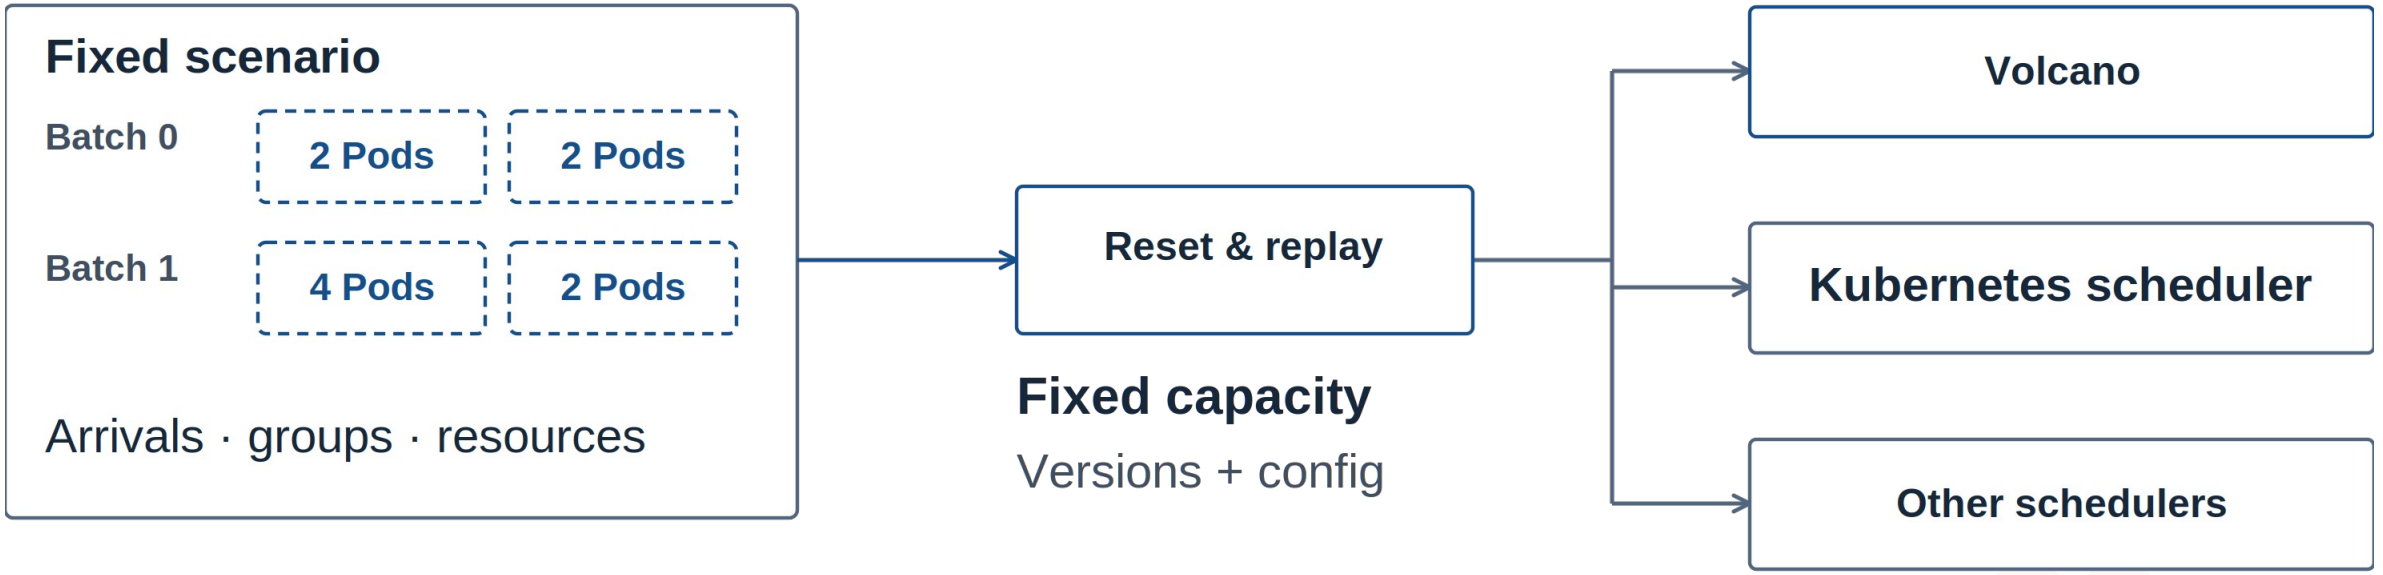
**04 Benchmark framework & progress**

# A Repeatable Scheduling Benchmark Framework



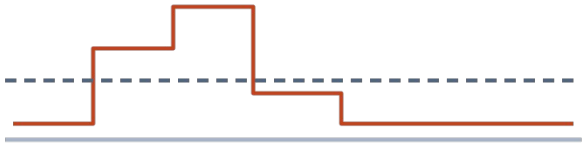
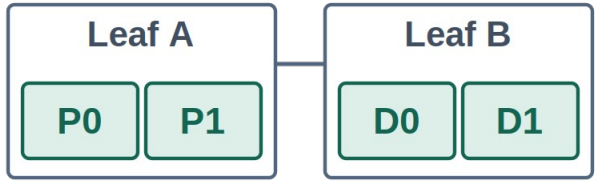
# Compare Schedulers on the Same Workload

## Planned comparisons



# Which Scenarios Fit Which Scheduler?

## Planned test matrix

| Scenario | | Variables | Metrics |
|---|---|---|---|
| **Bursts** |  | **Arrival rate · capacity**<br>Pool capacity requests | **Tail latency**<br>Backlog recovery |
| **Gang size** |  | **Group size · fragmentation**<br>Large / small gangs | **Complete gangs**<br>Binding throughput |
| **Topology** |  | **Tier limits · domains**<br>Gang · anti-affinity | **Valid placement**<br>Gang latency |

# What We Will Work On Next

| Workstream | Next | Validation |
| --- | --- | --- |
| **Group placement** | **SubGroup relationships**<br>Affinity / anti-affinity | Locality<br>Replica isolation |
| **Ecosystem** | **RBG · Kubeflow · LWS**<br>Multi-group integration | Group identity<br>Placement intent |
| **Benchmarks** | **Cross-scheduler testing**<br>Best-fit scenarios · topology | Latency · throughput<br>Placement correctness |

# Volcano community



# Thank you



**volcano.sh**



**GitHub**



**Slack**

Questions?